

COURSE OVERVIEW

CSC 105 - DIGITAL HUMANITIES - SPRINT 23

Marcus Birkenkrahe

January 6, 2023



What are "Digital Humanities" about?

- What do you think "Digital Humanities" are about?
- Compare: TOC of The Digital Humanities Cookbook (Drucker, 2021)
- Compare: Curriculum suggestion by ChatGPT (01-06-2023)




Figure 1: Map of West Africa in the Catalan Atlas (1375)

What will we do in this course?

- We're going to work with text data only (no maps¹)
- You can complete the "Text mining with R" track
- You will complete one DataCamp chapter per week
- In class, we will prepare and review the online material
- You will cover some of the most important AI/ML applications

How will you be evaluated?

- All course requirements have deadlines
- Late submissions will be penalized (loss of points)

¹Maps and GIS (Geographic Information Systems) are super cool but deserve an introductory course on their own. The focus of my teaching is to get you practical experience that you can use on the job instead of a broad conceptual overview (though the latter would be easier for all of us).

WEEK	DATE	DATA CAMP ASSIGNMENT
1	Jan 10,12	Wrangling text
2	Jan 17,19	Visualizing text
3	Jan 24,26	Sentiment analysis
4	Jan 31, Feb 2	String basics
5	Feb 7,9	Introduction to <code>stringr</code>
6	Feb 14,16	Pattern matching with regexp
7	Feb 21,23	More advanced matching
8	Mar 2	Three case studies
9	Mar 7,9	Text mining with Bag-of-Words
10	Mar 14,16	Word clouds and other visuals
11	Mar 28,30	Word clustering & tokenization
12	Apr 4,6	HR analytics case study
13	Apr 11,13	Polarity scoring
14	Apr 18,20	Visualizing sentiment
15	Apr 25,27	Case study: Airbnb reviews
16	May 2	

Figure 2: Source: syllabus, Canvas (lyon.instructure.com) or GitHub (github.com/birkenkrahe/ml)

REQUIREMENT	UNITS	PPU	TOTAL	% of TOTAL
Final exam	0	0	0	0
Home assignments	15	10	150	30.
Class assignments	10	10	100	20.
Final project	1	150	150	30.
Multiple-choice tests	10	10	100	20.
TOTAL			500	100.

Figure 3: Source: syllabus, Canvas (lyon.instructure.com) or GitHub (github.com/birkenkrahe/ml)

- No final exam! But weekly tests are graded
- The project topic can come from any of the course sub-topics
- The project deliverable is an essay of at least 5,000 words

Which tools are you going to use?

- DataCamp courses (15 weekly home assignments)
- GitHub repository (all course materials except tests)
- GNU Emacs + ESS + R (literate programming environment)
- Canvas (learning management system)

How can you register at DataCamp?

- You find the invitation link to the group for Spring 23 in Canvas.
- You will automatically be subscribed to the Digital Humanities team



Figure 4: Unsplash

- If you are in more than one course, I will add you later manually
- These accounts will be valid until July 8, 2023 only




	Introduction to Text Analysis in R Wrangling Text Chapter	Team	Active	Jan 17, 14:30 CST
	Introduction to Text Analysis in R Visualizing Text Chapter	Team	Active	Jan 24, 14:30 CST
	Introduction to Text Analysis in R Sentiment Analysis Chapter	Team	Active	Jan 31, 14:30 CST

Figure 5: DataCamp assignments for January

When is the first assignment due?



- The first DataCamp home assignment is due on January 17. For late submissions, you lose 1 point per day (out of 10 possible points)
- We'll write the first weekly multiple-choice test on January 19.

What should we do as a project?



Figure 6: Source: learning.edanz.com

- The **final essay** should be about one of the areas of Digital Humanities that we have **not** covered in our text mining course

- Here is the set of available essay topics (generated by an AI)
- An interesting approach would be if everyone picked the same topic related to text mining, ChatGPT being an obvious current "hot topic", and investigated different aspects, e.g. technical, ethical, societal, and personal aspects of the chatbot or this type of bot.
- You need to do your own research, including a literature review, and adhere to the IMRaD framework (see video):
 1. Introduction (what did you want to research?) with abstract
 2. Method (what did you do?) with literature review
 3. Results (what did you find out?) with examples, illustrations
 4. Discussion (what does it mean?) with limitations and outlook

What else could you do for a good start?

1. Complete/review introductory R or statistics courses:
 - Introduction to R" in DataCamp (data structures)
 - Introduction to statistics
 - fasteR by Norman Matloff (GitHub) - fast lane to R

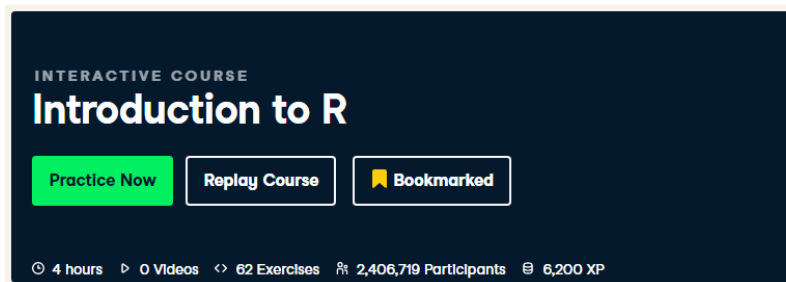


Figure 7: DataCamp course dashboard banner

2. If you do not have any experience with Emacs, work through the **online tutorial** (open it in Emacs with `CTRL + h t`) - ca. 1 hour.
 - Learn to open/close the editor
 - Learn basic cursor control (moving around)
 - Learn basic file management (open/close/find/save files)
 - Learn basic windows (buffer) management

What are you looking forward to?



- Learning more about text mining
- Unlocking the secrets of natural language processing
- Having fun with R programming
- Helping you on your own "digital humanities" journey

Next topic




Figure 8: Cover of Text Mining In Practice With R by Ted Kwartler (Wiley, 2010)