

Text mining in practice - Bag of Words - Intro to word clouds

Digital Humanities DSC 105 Spring 2023

Marcus Birkenkrahe

March 30, 2023

README

- This lecture closely follows the DataCamp lesson "Text Mining with Bag-of-Words in R" by Ted Kwartler, chapter 2, lesson 2, "Intro to word clouds" ([Link](#)).
- Download and open the practice file `8_wordclouds_practice.org` from GitHub to code along.
- In this lecture & practice:
 1. simple word cloud creation using Chardonnay tweets
 2. adding stop words to change word cloud
 3. improve word cloud appearance and insights

Get the corpus data and the R packages

- Download `corpora.R` from GitHub (bit.ly/tm-corpora)
- Run the file on the shell (`M-x eshell`) as a batch job:

```
R CMD BATCH corpora.R  
ls -al .RData
```

- Load the `.RData` file in your current R session and check that packages and user-defined objects were loaded:

```

load_packages <- function() {
  library(tm)
  library(qdap)
  library(SnowballC)
  library(wordcloud)
  search()
}
load_packages()
load("c:/Users/birkenkrahe/Downloads/.RData")
search()
ls()

[1] ".GlobalEnv"           "package:viridisLite"
[3] "package:wordcloud"    "package:SnowballC"
[5] "package:qdap"         "package:RColorBrewer"
[7] "package:qdapTools"    "package:qdapRegex"
[9] "package:qdapDictionaries" "package:tm"
[11] "package:NLP"          "ESSR"
[13] "package:stats"        "package:graphics"
[15] "package:grDevices"    "package:utils"
[17] "package:datasets"     "package:stringr"
[19] "package:httr"         "package:methods"
[21] "Autoloads"            "package:base"
[1] ".GlobalEnv"           "package:viridisLite"
[3] "package:wordcloud"    "package:SnowballC"
[5] "package:qdap"         "package:RColorBrewer"
[7] "package:qdapTools"    "package:qdapRegex"
[9] "package:qdapDictionaries" "package:tm"
[11] "package:NLP"          "ESSR"
[13] "package:stats"        "package:graphics"
[15] "package:grDevices"    "package:utils"
[17] "package:datasets"     "package:stringr"
[19] "package:httr"         "package:methods"
[21] "Autoloads"            "package:base"
[1] "api_key"              "ask_chatgpt"
[3] "chardonnay_corpus"    "chardonnay_df"
[5] "chardonnay_m"         "chardonnay_src"
[7] "chardonnay_tdm"       "chardonnay_vec"
[9] "clean_chardonnay"     "clean_chardonnay_corpus"
[11] "clean_coffee"         "clean_coffee_corpus"

```

```

[13] "coffee_corpus"      "coffee_df"
[15] "coffee_m"           "coffee_src"
[17] "coffee_tdm"         "coffee_vec"
[19] "color_pal"           "idx"
[21] "load_packages"       "m"
[23] "M"                   "max"
[25] "stops"               "term_frequency"
[27] "terms"               "word_freq"

```

- You need the `clean_coffee` and `clean_chardonnay` corpora.
- If we don't finish with a session, save your data from now on:

```

save.image(file=".RData")
shell("ls -al .RData")

```

```

-rwx-----+ 1 Birkenkrahe LYONNET+Group(513) 755072 Mar 30 08:53 .RData

```

Intro to word clouds

- Our starting point is the cleaned corpus of "Chardonnay" tweets, `clean_chardonnay`.
- Look at the corpus:

```

clean_chardonnay

<<VCorpus>>
Metadata: corpus specific: 0, document level (indexed): 0
Content: documents: 1000

```

Convert TDM to matrix

- Convert Chardonnay TDM to a matrix and check its dimensions:

```

library(tm)
chardonnay_tdm <- TermDocumentMatrix(clean_chardonnay)
chardonnay_m <- as.matrix(chardonnay_tdm)
dim(chardonnay_m)

[1] 3067 1000

```

Create the frequency table

- The starting point of any visualization is a frequency table - it only has two columns: terms (**term**) and their counts (**num**).
- Sum rows and sort by frequency:

```
term_frequency <- rowSums(chardonnay_m)
term_frequency <- sort(term_frequency, decreasing=TRUE)
word_freq <- data.frame(term = names(term_frequency),
                        num = term_frequency)
rownames(word_freq) <- NULL
head(word_freq, n=10)
```

	term	num
1	chardonnay	822
2	amp	120
3	marvin	104
4	wine	83
5	gaye	76
6	just	75
7	glass	63
8	like	55
9	bottle	47
10	lol	43

Add stop words and re-run the cleaning code

- The words **amp**, **wine** and **glass** do not help much - how can we get rid of them at this stage of our investigation? Do you know what "amp" means in this context?¹

Answer:

1. download the latest version of **corpora.R** from GitHub.
2. add these words to the stopwords cleaning function in **corpora.R**

¹Funnily enough, I had no idea until I looked into the raw CSV file: **amp** is a remnant of **&** after **removePunctuation**, and it's the HTML short code for **&**, which is frequent in tweets (saves 2 letters). As an interesting aside: I am already so dependent on ChatGPT that instead of checking the data, I went and asked the bot about "amp in the context of Chardonnay" but to no avail, of course.

3. run the batch job with R CMD BATCH

4. re-load .RData in this file.

You'll see that the number of words (records) has gone down and the list of top frequency words is changed.

- After cleaning out the additional words, reload the data, create the TDM and the word frequency data frame:

```
load("~/Downloads/.RData")
library(tm)
chardonnay_tdm <- TermDocumentMatrix(clean_chardonnay)
chardonnay_m <- as.matrix(chardonnay_tdm)
dim(chardonnay_m)
term_frequency <- rowSums(chardonnay_m)
term_frequency <- sort(term_frequency, decreasing=TRUE)
word_freq <- data.frame(term = names(term_frequency),
                        num = term_frequency)
rownames(word_freq) <- NULL
head(word_freq,n=10)
```

```
[1] 3067 1000
  term num
1  chardonnay 822
2      amp 120
3   marvin 104
4     wine  83
5    gaye  76
6    just  75
7   glass  63
8    like  55
9  bottle  47
10    lol  43
```

Using the wordcloud function

- We want to create word clouds. Is there a `wordcloud` function in `tm` or `qdap` or `base`? How can you find out? Load these packages (again, just in case) and check each of them for the function:

```
library(tm)
library(qdap)
library(wordcloud)
any(ls('package:tm')== "wordcloud")
any(ls('package:qdap')== "wordcloud")
any(ls('package:wordcloud')== "wordcloud")
```

```
[1] FALSE
[1] FALSE
[1] TRUE
```

- To create a wordcloud, use the `wordcloud` function. Look at the `help`.
- Use the column vectors `term` and `num` for the `words` and `freq` parameters, respectively:

```
library(wordcloud)
wordcloud(words=word_freq$term,
          freq=word_freq$num,
          max.words=100,
          color="blue")
```

tasting gaye http://duodifylink.com/moonson
 milkshake, bass wanna come composition inspired
 marvin thinkin' asse will be remember pretty really
 buzzes needin' slowin' down don't drinkin' call
 wine get drinkin' call don't drinkin' call
 happy way, richen cabernet best
 just price or you're glass glass best
 oaked full well name me get better charles
 full little night day thanks one's portrait
 meeting think make winnin' me love always
 white double winnin' me love always
 https double winnin' me love always
 way, richen cabernet best
 sip, lannuanyames lol like amp
 beauty wonderwines chipotlewines dant
 eduabudedubu

- ```
term_frequency[1:10]
```

Extract the terms 2 to 11 using `names` on `term_frequency` and call the vector of strings `terms_vec`. Show the entries 2 to 11:

7





```
content(clean_chardonnay[[24]])
```

```
[1] " brought marvin gaye chardonnay"
```

- You can add to the stopwords, and run `tm_map` with `removeWords` on the clean corpus to remove additional words:

```
content(clean_chardonnay[[24]])
stops <- c(stopwords("en"), 'just', 'like')
tail(stops)
clean_chardonnay_corpus <- tm_map(clean_chardonnay,
 removeWords,
 stops)
content(clean_chardonnay_corpus[[24]])
```

```
[1] " brought marvin gaye chardonnay"
[1] "so" "than" "too" "very" "just" "like"
[1] " brought marvin gaye chardonnay"
```

- To see the updated word cloud, re-run the code chunks from before with the new, cleaner corpus, then go back and rerun the last plot:

```
clean_chardonnay <- clean_chardonnay_corpus
library(tm)
chardonnay_tdm <- TermDocumentMatrix(clean_chardonnay)
chardonnay_m <- as.matrix(chardonnay_tdm)
dim(chardonnay_m)
term_frequency <- rowSums(chardonnay_m)
term_frequency <- sort(term_frequency, decreasing=TRUE)
word_freq <- data.frame(term = names(term_frequency),
 num = term_frequency)
rownames(word_freq) <- NULL
head(word_freq, n=10)
```

```
[1] 3065 1000
 term num
1 chardonnay 822
2 amp 120
3 marvin 104
```

```

4 wine 83
5 gaye 76
6 glass 63
7 bottle 47
8 lol 43
9 little 35
10 rose 34

```

## Improve word clouds with different colors

- The available colors are stored in a `character` vector `colors()`. Look at the `head` of the vector, and verify that 657 colors are available:

```

head(colors())
length(colors())

[1] "white" "aliceblue" "antiquewhite" "antiquewhite1"
[5] "antiquewhite2" "antiquewhite3"
[1] 657

```

- Instead of coloring all words with the same color, you can assign a vector of different colors to `wordcloud` to make certain words stand out or to fit a specific color scheme (e.g. to accommodate color-blind people).
- If you look at the `wordcloud` arguments:

```

load_packages <- function() {
 library(tm)
 library(qdap)
 library(SnowballC)
 library(wordcloud)
 search()
}
load_packages()
args(wordcloud)

[1] ".GlobalEnv" "package:viridisLite"
[3] "package:wordcloud" "package:SnowballC"

```

```

[5] "package:qdap" "package:RColorBrewer"
[7] "package:qdapTools" "package:qdapRegex"
[9] "package:qdapDictionaries" "package:tm"
[11] "package:NLP" "ESSR"
[13] "package:stats" "package:graphics"
[15] "package:grDevices" "package:utils"
[17] "package:datasets" "package:stringr"
[19] "package:httr" "package:methods"
[21] "Autoloads" "package:base"
function (words, freq, scale = c(4, 0.5), min.freq = 3, max.words = Inf,
 random.order = TRUE, random.color = FALSE, rot.per = 0.1,
 colors = "black", ordered.colors = FALSE, use.r.layout = FALSE,
 fixed.asp = TRUE, ...)
NULL

```

- The `colors` argument colors words from least to most frequent. The code uses three colors of increasing vibrancy - this will naturally divide the term frequency into "low", "medium", and "high":

```

term_frequency <- rowSums(chardonnay_m)
term_frequency <- sort(term_frequency, decreasing=TRUE)
word_freq <- data.frame(term = names(term_frequency),
 num = term_frequency)
rownames(word_freq) <- NULL
head(word_freq,n=10)
wordcloud(words=word_freq$term,
 freq=word_freq$num,
 max.words=100,
 colors=c("grey80","darkgoldenrod1","tomato"))

```



```

[1] ".GlobalEnv" "package:viridisLite"
[3] "package:wordcloud" "package:SnowballC"
[5] "package:qdap" "package:RColorBrewer"
[7] "package:qdapTools" "package:qdapRegex"
[9] "package:qdapDictionaries" "package:tm"
[11] "package:NLP" "ESSR"
[13] "package:stats" "package:graphics"
[15] "package:grDevices" "package:utils"
[17] "package:datasets" "package:stringr"
[19] "package:httr" "package:methods"
[21] "Autoloads" "package:base"

```

- Look at the contents of the package with `ls`: these are the different color maps.

```
ls('package:viridisLite')
```

```

[1] "cividis" "inferno" "magma" "mako" "plasma"
[6] "rocket" "turbo" "viridis" "viridis.map" "viridisMap"

```

- All maps are functions with one mandatory argument, the number of colors `n` used. Check the arguments of `viridisLite::cividis`:

```
args(cividis)
```

```
function (n, alpha = 1, begin = 0, end = 1, direction = 1)
NULL
```

- As the vignette for `viridisLite` reveals, the other parameter allow to change transparency (`alpha`), hue (`begin` and `end`), and order. Here are the color scales for the maps:



- To create a new wordcloud with the selected palette, select 5 colors from `turbo` and store them in a vector `color_pal`:

```
color_pal <- turbo(5)
```

- Print the hex-codes for `color_pal` to the console:

```
color_pal
```

```
[1] "#30123BFF" "#28BBECFF" "#A2FC3CFF" "#FB8022FF" "#7A0403FF"
```

- Create a word cloud from the Chardonnay tweets `word_freq`, include 100 terms, and set the `colors` to the `color_pal` palette:

```
wordcloud(words=word_freq$term,
 freq=word_freq$num,
 max.words=100,
 colors=color_pal)
```

chardonnay



•

## Load packages

```
load_packages <- function() {
 library(tm)
```

```

 library(qdap)
 library(SnowballC)
 library(wordcloud)
 search()
}
load_packages()

[1] ".GlobalEnv" "package:viridisLite"
[3] "package:wordcloud" "package:SnowballC"
[5] "package:qdap" "package:RColorBrewer"
[7] "package:qdapTools" "package:qdapRegex"
[9] "package:qdapDictionaries" "package:tm"
[11] "package:NLP" "ESSR"
[13] "package:stats" "package:graphics"
[15] "package:grDevices" "package:utils"
[17] "package:datasets" "package:stringr"
[19] "package:httr" "package:methods"
[21] "Autoloads" "package:base"

```